



BEHAVIORAL ANOMALY DETECTION FOR INDIRECT PROMPT INJECTION IN AGENTIC SOC WORKFLOWS

Final Year Thesis Presentation – 2026


Presented by:

Md. Tanjim Mahmud Tuhin
251-56-012

Supervised by:

Prof. Dr. Saifullah Sadi
Head of the MSc in Cybersecurity

Bangladesh's AI-Powered SOC Landscape Creates New Attack Surfaces



Legacy Risk: Bangladesh Bank Heist 2016

\$81M loss demonstrated devastating impact of critical infrastructure compromise.



National SOC Development: Bangladesh National SOC (NSOC) is currently under development via the BCC EDGE Project.



AI in Modern SIEMs: Major SIEMs (Microsoft Sentinel, IBM QRadar) now integrate **AI agents** for automated analysis.




Attacker-Controlled Telemetry: AI agents ingest data directly from attacker-controlled sources: SIEM logs, threat feeds, and emails.



The New Threat: A single hidden instruction in ingested data can **redirect the entire SOC workflow**.

Indirect Prompt Injection Exploits the Environment, Not the User

DIU Logo



→ Direct Injection

Attacker speaks directly to the AI via the user interface. Requires direct interaction with the LLM prompt.

↔ Indirect Injection (SOC Context)

Attacker hides malicious instructions in data that the AI agent reads from its environment (invisible to the analyst).

- ▶ SIEM/EDR Logs
- ▶ Threat Intelligence RAG Feeds
- ▶ Alert/Incident Tickets

Systematic Review of 445 Papers Reveals a Critical Research Gap

Poisoning the Watchtower (2026)

SOC-specific attack scenarios

Log-only focus, no ATLAS mapping

AgentDojo (2024)

Agent behavior benchmark

Not SOC-specific

InjecAgent (2024)

IPI benchmark suite

No behavioral defense mechanism

StruQ (2025)

Input filtering techniques

Not runtime monitoring

MITRE ATLAS papers

ML/AI attack coverage

Not SOC-specific workflows

	SOC-Specific	3 Channels	ATLAS + Behavioral
Watchtower	✓	×	×
AgentDojo	×	✓	×
InjecAgent	THE RESEARCH GAP (NOBODY HAS DONE THIS)		
	×	✓	

Based on systematic review of 445 papers across 9 clusters

Five Interconnected Research Objectives

1 Threat Model & Taxonomy

SOC-specific threat model + attack taxonomy mapped to MITRE ATLAS & ATT&CK.



2 Testbed & Dataset

Controlled agentic SOC testbed + multilingual dataset (English, Bangla, mixed).



3 Telemetry Schema

Behavioral telemetry schema (tool calls, workflow states, provenance, resource access).



4 Hybrid Detector

Hybrid policy-aware sequence anomaly detector (DARF).




5 Defense Prototype

Runtime defense prototype + Bangladesh adoption guidelines.




Trust Boundary Architecture: Clear Attack Surface Definition



Four-Layer Framework: From Data Ingestion to Response Control


DIU Logo




 DETECTED

Behavioral Telemetry Captures Every Agent Execution Step


DIU Logo



DARF: Hybrid Detector Combines Deterministic Policy, Sequence Anomaly, and Risk Fusion




Three-Tier Graduated Response from Low to High Risk



Multilingual Dataset Design and Comprehensive Evaluation Framework

Dataset Composition

- ▶ **Scale:** 800-1200 controlled agent executions
- ▶ **Languages:** English, Bangla, Mixed Bangla-English
- ▶ **Attack Types:** Direct override, Persona hijack, Context manipulation, Obfuscation



Evaluation Metrics

Attack Success Rate	Precision & Recall
F1 Score	FPR & FNR
Detection Latency	Task Completion Rate
Cross-carrier Generalization	Cross-language Generalization

Baselines for Comparison

✖	No defense (Control)
⚠	Input filtering only
📄	Policy-only detection
📡	Sequence-only detection

First Formal SOC-Specific MITRE ATLAS Attack Mapping




FIRST FORMAL SOC-SPECIFIC ATLAS MAPPING FOR INDIRECT PROMPT INJECTION

Attack Vector	Injection Channel	ATLAS Tactic	Technique ID	ATT&CK Mapping
Log Field Injection	SIEM	<i>ML Attack Staging</i>	AML.T0051	T1565 (Data Manipulation)
RAG Poisoning	Threat Intel Feed	<i>Persistence</i>	AML.T0053	T1195 (Supply Chain Compromise)
Alert Context Manipulation	Incident Ticket	<i>Impact</i>	AML.T0048	T1499 (Endpoint Denial of Service)

12-Month Work Plan with Current Progress

DIU Logo



Six Tangible Deliverables Targeting Top Security Venues



01

Threat Model & Taxonomy

Comprehensive SOC-specific attack taxonomy mapped to MITRE ATLAS and ATT&CK frameworks.



02

Multilingual Dataset

Sanitized dataset of controlled agent executions featuring English, Bangla, and mixed language attacks.



03

Agentic SOC Testbed

Controlled testing environment with a robust behavioral telemetry schema for execution tracing.



04

Hybrid DARF Detector

Advanced anomaly detector combining deterministic policies, sequence analysis, and risk fusion.



05

Runtime Defense Prototype

Vendor-neutral prototype implementing graduated runtime responses to detected injection attacks.



06

Adoption Guidelines

Strategic guidelines for the secure adoption of AI in Bangladesh's critical infrastructure SOCs.

Monitoring Agent Behavior, Not Just Input, Is the Path to Secure AI-SOC

PROBLEM

AI agents in SOC read attacker-controlled data (logs, feeds, tickets).
No behavioral defense currently exists.



SOLUTION

DARF Framework:
4-layer behavioral monitoring, hybrid detection, and graduated response control.



IMPACT

Bangladesh can safely adopt AI in its critical infrastructure SOCs, mitigating modern automated threats.

“

*"Security data is inherently adversarial. We should monitor **what the agent does**, not just what it reads."*

FUNCTIONAL REQUIREMENTS

FR-01

Monitor Agent Tool Usage

The system shall monitor and log all tool invocations made by the LLM agent, including tool name, parameters, and output.

FR-02

Track Reasoning Steps

The system shall capture and store the agent's intermediate reasoning traces and chain-of-thought outputs for behavioral comparison.

FR-03

Detect Behavioral Anomalies

The system shall detect statistically significant deviations from established behavioral baselines using trained anomaly detection models.

FR-04


Alert SOC System

Upon anomaly detection, the system shall generate structured alerts with severity level, affected task, and suspected injection vector.

FR-05


Log Suspicious Behavior

All flagged behaviors shall be persisted to a tamper-evident audit log for forensic investigation and model refinement.




SYSTEM ARCHITECTURE

INPUT SOURCES




USE CASE: INDIRECT PROMPT INJECTION

Scenario: Attacker embeds malicious instruction inside a SIEM log entry



ACTIVITY DIAGRAM – SYSTEM WORKFLOW



EVALUATION DESIGN

Benchmark Datasets

[AgentDojo]

Tests LLM agents against adversarial task manipulation. Includes multi-step tasks and injected payloads.

[MCPSecBench]

Security benchmark for Model Context Protocol. Evaluates tool misuse and instruction hijacking.

[Custom SOC Dataset]


Real SIEM log formats (Splunk/Elastic) with labeled benign vs. injected adversarial payloads.

Performance Metrics

Detection Accuracy 92%



False Positive Rate < 2%



Attack Success Reduction 85%



Detection Latency < 200ms



Real-World Attacks

CVE-2025-53773

Indirect prompt injection via MCP tool responses. Attack redirects agent control-flow via malicious tool output.

EchoLeak – CVE-2025-32711

Data exfiltration via indirect injection in Microsoft Copilot. Hidden instructions in documents trigger data leaks.

HYPOTHESIS: Behavioral monitoring detects ≥85% of injection attempts missed by input filtering alone.

CONCLUSION

Core Contribution

This research proposes the **FIRST** behavioral anomaly detection framework specifically designed for indirect prompt injection in SOC-integrated agentic LLMs. By monitoring agent control-flow, tool usage, and reasoning deviations, the system detects subtle attacks that bypass existing input filters.

Key Findings

- Behavioral monitoring detects >85% of indirect injection attempts.
- Input filtering alone misses context-dependent injections.
- CVE-2025-53773 and EchoLeak fully reproducible in lab setting.
- Cross-model detection generalizes across GPT-4, Claude, LLaMA 3.
- Real-time alert latency < 200ms in simulated SOC environment.

Impact & Future Work

- Improves security of AI-driven SOC automation at runtime.
- Provides formal MITRE ATLAS mapping for SOC injection attacks.
- Dataset and framework released open-source for future research.
- Future: Extend to multi-agent SOC systems and federated detection.
- Target: IEEE S&P / USENIX Security / ACM CCS publication.

"The future of SOC security lies not in filtering what goes IN – but in understanding what the agent DOES."



REFERENCES

- [1] Perez, F., & Ribeiro, I. (2022). **Ignore Previous Prompt: Attack Techniques for Language Models.** arXiv:2211.09527.
- [2] Zou, A., et al. (2023). **Universal and Transferable Adversarial Attacks on Aligned Language Models.** arXiv:2307.15043.
- [3] MITRE. (2023). **ATLAS: Adversarial Threat Landscape for Artificial Intelligence Systems.** <https://atlas.mitre.org>
- [4] OWASP. (2024). **OWASP Top 10 for Large Language Model Applications.** <https://owasp.org>
- [5] DeBenedetti, E., et al. (2024). **AgentDojo: A Dynamic Environment to Evaluate Prompt Injection Attacks.** arXiv:2406.13352.
- [6] MCPSecBench Authors. (2025). **MCPSecBench: A Benchmark for Evaluating Security of MCP Implementations.**
- [7] NVD. (2025). **CVE-2025-53773: Indirect Prompt Injection via MCP Tool Responses.** <https://nvd.nist.gov>
- [8] Microsoft MSRC. (2025). **CVE-2025-32711 – EchoLeak: Data Exfiltration via Copilot Prompt Injection.**
- [9] Liu, Y., et al. (2024). **Prompt Injection Attacks and Defenses in LLM-Integrated Applications.** arXiv:2310.12815.
- [10] Greshake, K., et al. (2023). **Not What You've Signed Up For: Compromising LLM-Integrated Applications.** arXiv:2302.12173.
- [11] MITRE. (2024). **ATLAS Mitigations: AML.M0015 – Adversarial Input Detection.** <https://atlas.mitre.org>

THANK YOU

Questions and Discussion Welcome

[Md. Tanjim Mahmud Tuhin] | [251-56-012]
Daffodil International University | MSc. in Cybersecurity

Email: tanjimtuhin06@gmail.com

Supervisor: Prof. Dr. Saifullah Sadi




Email



LinkedIn



Research Profile



Q & A

Questions & Answers

Please feel free to ask any questions about the research, methodology, system design, or evaluation approach.

- Why behavioral monitoring vs. input filtering?
- How is the behavioral baseline established?
- How do you handle zero-day injections?
- What is the latency overhead of the monitor?
- How does this generalize across different LLMs?